

10 Week 10: Machine Learning for Injury Risk

This code sets up, trains, and evaluates two machine learning models (Logistic Regression and Random Forest) to predict injury risk.

```
# Week 10: Machine Learning for Injury Risk Prediction
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix

# --- Data Loading (Assumes 'bulls_injury_data.csv' was generated in the setup cell) ---
df = pd.read_csv('bulls_injury_data.csv')

# --- 1. Define Features (X) and Target (y) ---
X = df[['Minutes_Load', 'HighSpeedRuns', 'JumpLoad', 'HeartRate', 'SleepHours', 'PreviousInjury']]
y = df['InjuryNextWeek']

# --- 2. Split the Data (80% Train, 20% Test) ---
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
```

```
)
```

```
# --- 3. Scale Features (Essential for Logistic Regression) ---
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# --- 4. Train and Evaluate Logistic Regression ---
```

```
print("\n--- Logistic Regression ---")
```

```
log_reg = LogisticRegression(random_state=42)
```

```
log_reg.fit(X_train_scaled, y_train)
```

```
y_pred_log = log_reg.predict(X_test_scaled)
```

```
print("📊 Classification Report:")
```

```
print(classification_report(y_test, y_pred_log, zero_division=0)) # zero_division=0 handles  
the imbalance
```

```
# --- 5. Train and Evaluate Random Forest ---
```

```
print("\n--- Random Forest ---")
```

```
# Using original X data (no scaling needed for tree-based models)
```

```
rf = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
rf.fit(X_train, y_train)
```

```
y_pred_rf = rf.predict(X_test)
```

```
print("📊 Classification Report:")
```

```
print(classification_report(y_test, y_pred_rf, zero_division=0))
```

```
# --- 6. Visualize Feature Importance ---
```

```
feat_importance = pd.Series(rf.feature_importances_, index=X.columns)
```

```
plt.figure(figsize=(8,5))
```

```
feat_importance.sort_values().plot(kind='barh', color='red')
```

```
plt.title('Feature Importance - Injury Risk Prediction (Bulls)')
```

```
plt.show()
```

```
# --- 7. Confusion Matrix (Final Visualization) ---
```

```
cm = confusion_matrix(y_test, y_pred_rf, labels=[0, 1]) # Ensure correct labels are included
```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Reds', xticklabels=['No Injury', 'Injury'],  
            yticklabels=['No Injury', 'Injury'])
```

```
plt.title('Confusion Matrix - Random Forest')
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.show()
```